

UQ score vs selection stringency (top-k%) — per-model mut cap

left = most selective slice; k = 100% = full-set mean

